

PROFILE

Software Engineer (UNSAAC, graduated coursework) specialized in building AI systems end to end: production agents with RAG, tool-calling and full traceability, backed by active study of LLM internals (attention, RoPE, SwiGLU) to debug them with real technical judgment. Hands-on experience deploying AI pipelines on AWS Lambda and Vercel. Looking for a Junior AI Engineer role to take LLM prototypes into reliable systems.

TECHNICAL SKILLS

Python

PyTorch

Local LLMs (Ollama)

PydanticAI

RAG

Qdrant

LanceDB / GraphRAG

Prompt Engineering

Fine-tuning / LoRA

FastAPI

OpenTelemetry

Next.js

AWS Lambda

Docker

SQL

Git/GitHub

LANGUAGES

Spanish (Native)

English (B1, technical)

Quechua

EVENTS

NASA Space Apps Challenge Hackathon

2025

Official participant in NASA's global technology innovation event.

EDUCATION

B.Eng. in Informatics and Systems Engineering

UNSAAC

2021 - 2026

Cusco, Peru

Coursework complete; Bachelor's degree (Bachillerato) in progress. Focus on software development and applied Artificial Intelligence techniques.

EXPERIENCE

Research Assistant

LAAD, UNSAAC

2023 - 2025

Cusco, Peru

- Algorithms and Data Analysis Lab.
- Built the Next.js frontend for the "Conflicto de Tinta" historical RAG assistant, consuming a RAG backend deployed on AWS Lambda and presenting answers with verifiable sources.
 - Contributed to computer vision and NLP projects: OCR, data cleaning and labeling, and Python pipeline automation.
 - (2023) Built the lab's website with HTML, CSS and JavaScript.

Frontend & Project Management Intern (Hybrid)

Turix

Dec 2025 - Aug 2026

Cusco, Peru (Remote)

- Built Next.js and Astro frontends for real clients (Tour Operador SaaS, Magic Experiences Peru, Perou Magique Tours, Kusikuy Travel Transportes) in a remote, agile environment.
- Evolved into a hybrid role covering task coordination (Trello) and development, gathering technical requirements directly with clients.

FEATURED PROJECTS

Asistente Mikha (Personal AI Agent)

2026

Personal agent built with **FastAPI + PydanticAI** on a local LLM (Ollama): read-only tool-calling, explicit-confirmation actions, persistent memory over Markdown notes, and **OpenTelemetry** tracing to Phoenix for agent observability.

GraphRAG – Historical Assistant

2025

Built during my research at LAAD. RAG pipeline with **Microsoft GraphRAG + LanceDB** over Peruvian colonial-era archives; Next.js

Paqarina Wasi Incubator (UNSAAC)

Pre-incubation

📅 Mar 2026

Selected with **Yuyana**, a **fintech** for P2P loans and debt tracking, for the pre-incubation program at the Paqarina Wasi Incubator.

ADDITIONAL TRAINING

Studying LLM Architecture (PyTorch)

LangGraph

Deep Learning (CNNs)

frontend with citable sources and document preview. **Deployed on AWS Lambda + Vercel.**

Semantic Search Application 2025

Full-stack RAG app (**FastAPI + Qdrant + Next.js**) with dual embedding support (OpenAI and open-source) and a custom sentence-boundary-aware chunking pipeline.